

BioBERT Fine Tuning

Class 0 ("DLT") → Recall = 0.94

Class 1 ("Non-serious") → Recall = 1.00

Class 2 ("Serious") → Recall = 0.82

```
! pip install transformers
```

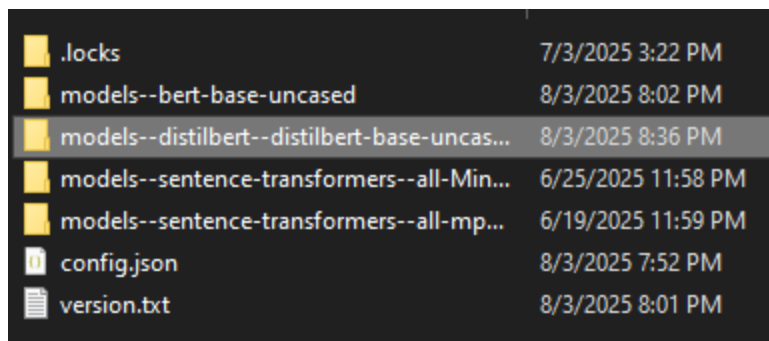
- BERT adds **CLS** at the beginning and **SEP** at end of the sentence.

```
from transformers import pipeline
```

- Pipeline helps you to call pretrained model

```
from transformers import pipeline  
classifier = pipeline('sentiment-analysis')
```

- This will download **models--distilbert--distilbert-base-uncased-finetuned-sst-2-english** model by default (255 MB)



.locks	7/3/2025 3:22 PM
models--bert-base-uncased	8/3/2025 8:02 PM
models--distilbert--distilbert-base-uncas...	8/3/2025 8:36 PM
models--sentence-transformers--all-Min...	6/25/2025 11:58 PM
models--sentence-transformers--all-mp...	6/19/2025 11:59 PM
config.json	8/3/2025 7:52 PM
version.txt	8/3/2025 8:01 PM

```
classifier("The movie is awesome")
```

```
[{'label': 'POSITIVE', 'score': 0.9998778104782104}]
```

Multiple test:

- Provide list

```
classifier(["We are very happy to show you the 😊 Transformers library.",  
           "We hope you don't hate it."])
```

```
[{'label': 'POSITIVE', 'score': 0.9997795224189758},  
 {'label': 'NEGATIVE', 'score': 0.5308613777160645}]
```

Fine Tuning using BioBERT ([dmis-lab/biobert-base-cased-v1.1](#))

📌 Official model name:

```
dmis-lab/biobert-base-cased-v1.1
```

Steps:

1. Call pretrained model
2. Tokenize using BioBERT
3. Convert the encodings into dataset objects



Load and Prepare Your CSV Data

```
import pandas as pd  
  
# Load your CSV file  
df = pd.read_csv("PV Classification.csv")  
df
```

	ID	Seriousness	Narrative
0	1	Non-serious	Two hours after taking a new antihistamine for...
1	2	D/LT	Within an hour of receiving a new chemotherapy...
2	3	Non-serious	On day three of a new antibiotic for a sinus i...
3	4	Serious	Five days after starting a new anticoagulant, ...
4	5	Non-serious	After my flu shot, I developed a low-grade fev...
...
645	646	Non-serious	My child developed mild diarrhea after receivi...
646	647	Serious	After administration of a long-acting injectab...
647	648	D/LT	Patient with chronic kidney disease received g...
648	649	Non-serious	Applying a eucalyptus balm caused a temporary ...
649	650	Serious	I started on isoniazid for latent TB and devel...

650 rows × 3 columns

Convert into list

```
# Only keep the text and labels
texts = df["Narrative"].tolist()
labels = df["Seriousness"].tolist()
```

```
texts
✓ 0.0s

['Two hours after taking a new antihistamine for seasonal allergies
'Within an hour of receiving a new chemotherapy infusion, I experi
'On day three of a new antibiotic for a sinus infection, I noticed
'Five days after starting a new anticoagulant, I experienced sever
'After my flu shot, I developed a low-grade fever of 100.1°F and m
'Two hours after taking a new painkiller for chronic migraines, I
'After applying a new topical steroid for eczema, I noticed mild i
'A week after starting a new diabetes medication, I noticed yellow
'On the second day of taking a new cough syrup, I experienced a dr
```

```
labels
✓ 0.0s
['Non-serious',
'D/LT',
'Non-serious',
'Serious',
'Non-serious',
'D/LT',
'Non-serious',
'Serious',
'Non-serious',
'D/LT',
'Non-serious']
```

Convert Labels to Numbers

```
from sklearn.preprocessing import LabelEncoder
```

```
label_encoder = LabelEncoder()
```

```
labels_encoded = label_encoder.fit_transform(labels)
```

```
labels_encoded
```

✓ 0.0s

```
array([[1, 0, 1, 2, 1, 0, 1, 2, 1, 0, 1, 2, 1, 0, 1, 2, 1, 0, 1, 2, 1, 0,  
       1, 2, 1, 0, 1, 2, 1, 0, 1, 2, 1, 0, 1, 2, 1, 0, 1, 2, 1, 0, 1, 2,  
       1, 0, 1, 2, 1, 0, 1, 0, 1, 2, 1, 0, 1, 2, 1, 0, 1, 2, 1, 0, 1, 2,  
       1, 0, 1, 2, 1, 0, 1, 2, 1, 0, 1, 2, 1, 0, 1, 2, 1, 0, 1, 2, 1, 0,  
       1, 2, 1, 0, 1, 2, 1, 0, 1, 2, 1, 0, 1, 0, 1, 2, 1, 0, 1, 2, 1, 0,  
       1, 2, 1, 0, 1, 2, 1, 0, 1, 2, 1, 0, 1, 2, 1, 0, 1, 2, 1, 0, 1, 2,  
       1, 0, 1, 2, 1, 0, 1, 2, 1, 0, 1, 2, 1, 0, 1, 2, 1, 0, 1, 0, 1, 2,  
       1, 0, 1, 2, 1, 0, 1, 2, 1, 0, 1, 2, 1, 0, 1, 2, 1, 0, 1, 2, 1, 0,  
       1, 2, 1, 0, 1, 2, 1, 0, 1, 2, 1, 0, 1, 2, 1, 0, 1, 2, 1, 0, 1, 2,
```

`LabelEncoder` converts text labels like "Non-serious", "D/LT" to integers like 0, 1, 2.

Load BioBERT Tokenizer

```
from transformers import AutoTokenizer
```

```
tokenizer = AutoTokenizer.from_pretrained("dmis-lab/biobert-base-cased-v1.1", cache_dir="hf_models")
```

This loads **BioBERT** tokenizer. It breaks the clinical sentence into special tokens that the model can understand.

Tokenize Text

```
# Convert text to input IDs and attention masks
encodings = tokenizer(texts, truncation=True, padding=True, max_length=256, return_tensors="tf")
```

- `truncation=True` : Cuts off long texts. (**Default Value:** `False` : truncation is **not applied** by default)
- `padding=True` : Ensures all sequences are of the same length. (**Default Value:** `False` (padding is **not applied** by default))
- `return_tensors="tf"` : Converts to TensorFlow tensors.

```
encodings
✓ 0.0s

{'input_ids': <tf.Tensor: shape=(650, 103), dtype=int32, numpy=
array([[ 101, 1160, 2005, ..., 0, 0, 0],
       [ 101, 1439, 1126, ..., 8182, 119, 102],
       [ 101, 1113, 1285, ..., 0, 0, 0],
       ...,
       [ 101, 5351, 1114, ..., 0, 0, 0],
       [ 101, 11892, 170, ..., 0, 0, 0],
       [ 101, 178, 1408, ..., 0, 0, 0]])>, 'token_type_ids': <tf.Tensor: shape=(650, 103), dtype=int32, numpy=
array([[0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       ...,
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0],
       [0, 0, 0, ..., 0, 0, 0]])>, 'attention_mask': <tf.Tensor: shape=(650, 103), dtype=int32, numpy=
array([[1, 1, 1, ..., 0, 0, 0],
       [1, 1, 1, ..., 1, 1, 1],
       [1, 1, 1, ..., 0, 0, 0],
       ...,
       [1, 1, 1, ..., 0, 0, 0],
       [1, 1, 1, ..., 0, 0, 0],
       [1, 1, 1, ..., 0, 0, 0]])>}
```

Prepare Dataset for Training

Convert to dataset:

```
import tensorflow as tf
from tensorflow.data import Dataset

dataset = Dataset.from_tensor_slices(
    (dict(encodings), labels_encoded)
)
```

Shuffle and split:

```
train_size = int(0.8 * len(dataset))
train_dataset = dataset.take(train_size).shuffle(100).batch(8)
val_dataset = dataset.skip(train_size).batch(8)
```

```
int(0.8 * len(dataset))
```

Output: 520

`dataset.take(train_size)` : take the first `train_size` samples from the dataset.

`.shuffle(100)` : The argument `100` is the **buffer size** used for shuffling. It means that it will randomly shuffle the dataset using a buffer of 100 elements at a time.

`dataset.batch(8)`

- **Purpose:** The `batch()` function groups samples into batches, which is crucial for efficient training.

- The argument `8` is the **batch size**, which means the dataset will be split into batches of 8 samples each.
- **Why batching?:** Training a model on the entire dataset at once might be very memory-intensive, especially for large datasets. Instead, you train your model on small chunks or **mini-batches** of data.
 - A batch size of 8 means the model will receive 8 samples at a time for training

Load BioBERT Model for Classification

```
from transformers import TFAutoModelForSequenceClassification

# Load the PyTorch weights and convert them to TensorFlow
model = TFAutoModelForSequenceClassification.from_pretrained(
    "dmis-lab/biobert-base-cased-v1.1",
    num_labels=3,
    from_pt=True, # This tells it to load from PyTorch weights
    cache_dir="hf_models"
)
```

- `TFAutoModelForSequenceClassification` : Loads BERT with classification head on top.
- `num_labels=3` : Number of output categories.

Compile the Model

```
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.losses import SparseCategoricalCrossentropy

optimizer = Adam(learning_rate=5e-5)
loss = SparseCategoricalCrossentropy(from_logits=True)

model.compile(optimizer=optimizer, loss=loss, metrics=["accuracy"])
```

Train the Model

```
model.fit(
    train_dataset,
    validation_data=val_dataset,
    epochs=5
)
```

```
Epoch 1/5
WARNING:tensorflow:From c:\Users\Jeevan\miniconda3\envs\bert_env\lib\site-packages\keras\src\utils\tf_utils.py:492: The name tf.ragged.RaggedTensor is deprecated. Please use tf.experimental.numpy.RaggedTensor instead.
WARNING:tensorflow:From c:\Users\Jeevan\miniconda3\envs\bert_env\lib\site-packages\keras\src\engine\base_layer_utils.py:384: The name tf.nn.conv2d is deprecated. Please use tf.nn.conv2d_v2 instead.

65/65 [=====] - 181s 2s/step - loss: 0.5529 - accuracy: 0.7596 - val_loss: 0.3012 - val_accuracy: 0.8538
Epoch 2/5
65/65 [=====] - 152s 2s/step - loss: 0.2486 - accuracy: 0.9077 - val_loss: 0.5364 - val_accuracy: 0.8308
Epoch 3/5
65/65 [=====] - 155s 2s/step - loss: 0.0893 - accuracy: 0.9673 - val_loss: 0.3624 - val_accuracy: 0.8923
Epoch 4/5
65/65 [=====] - 155s 2s/step - loss: 0.0448 - accuracy: 0.9904 - val_loss: 0.5139 - val_accuracy: 0.8538
Epoch 5/5
65/65 [=====] - 151s 2s/step - loss: 0.0400 - accuracy: 0.9827 - val_loss: 0.2268 - val_accuracy: 0.9308
<keras.src.callbacks.History at 0x1c4ed01ffa0>
```



Took 13+ Mins

Common question:

◆ 1. Was it necessary to convert **Narrative** and **Seriousness** into a list?

✓ Yes

- HuggingFace **Tokenizer** expects a **list of strings**, where each item is one **sentence** (narrative).
- Your data is in a **pandas DataFrame**. Each column is a **Series**, not a list.

◆ 2. What is **max_length=256** ?

Each sentence is converted into a maximum of 256 **tokens**.

If the sentence is longer → it gets **truncated**.

If it's shorter → it gets **padded** (extra `[PAD]` tokens added to reach 256 length).

Evaluate and Predict

◆ 3. Why we used `SparseCategoricalCrossentropy` as loss?

Because our labels are:

- **Categorical**, not continuous (e.g., "non-serious", "serious", "death")
- But **integer-encoded**, not one-hot encoded

You have 3 classes:

```
Non-serious → 0
D/LT → 1
Serious → 2
```

And you pass them as integers. So you can't use `CategoricalCrossentropy` (which expects one-hot labels like `[1, 0, 0]`).

◆ 4. What is `from_logits=True` ?

A **logit** is a raw output from the model **before applying softmax**.

Let's say:

```
model output → [1.5, 0.5, -2.0] (called logits)
```

To convert it into probabilities:

```
softmax([1.5, 0.5, -2.0]) → [0.73, 0.26, 0.01]
```

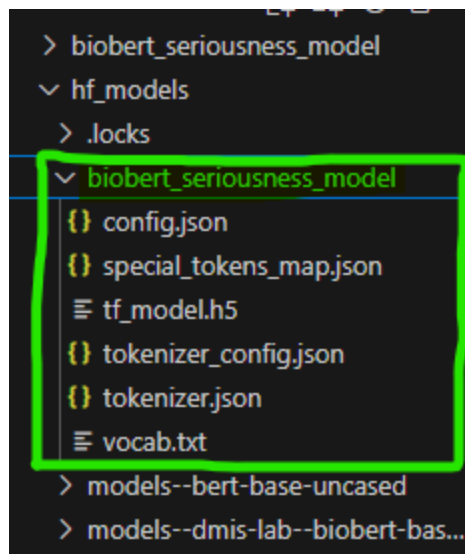
If you're using the model output directly, you should tell the loss function:

“Hey, I’m giving you logits, please apply softmax inside.”

That’s what `from_logits=True` does. If you set it to `False`, the model output must already be softmaxed.

Save the Model

```
model.save_pretrained("./hf_models/biobert_seriousness_model")
tokenizer.save_pretrained("./hf_models/biobert_seriousness_model")
```



✓ Load the Saved Model

```
from transformers import TFBertForSequenceClassification, BertTokenizer

# Load from the directory you saved to
model = TFBertForSequenceClassification.from_pretrained("./hf_models/biobert_seriousness_model")
tokenizer = BertTokenizer.from_pretrained("./hf_models/biobert_seriousness_model")
```

```

import tensorflow as tf

def predict_seriousness(text):
    # Tokenize text
    inputs = tokenizer(text, return_tensors="tf", truncation=True, max_length=128)

    # Get model predictions
    outputs = model(**inputs)

    # Convert logits to probabilities (for all 3 classes)
    proba = tf.nn.softmax(outputs.logits).numpy()[0]

    # Get predicted class (0, 1, or 2)
    predicted_class = tf.argmax(proba).numpy()

    # Map to label names
    label_names = ["DLT", "Non-serious", "Serious"] # Must match your training labels
    return label_names[predicted_class], "Confidence:", proba[predicted_class]

```

The screenshot shows a Jupyter Notebook interface with two code cells. The first cell contains the text tokenization code, and the second cell contains the model prediction code. Both cells show successful execution with timing information.

```

text = "Two weeks after starting a new heart medication, I experienced severe fatigue and fainting spells. I was admitted for suspected arrhythmia. The drug was stopped, and I received a pacemaker. I stabilized after five days."
✓ 0.0s Python

```

```

inputs = tokenizer(text, return_tensors="tf", truncation=True, max_length=128)
input
✓ 0.0s Python
<bound method Kernel.raw_input of <ipykernel.ipkernel.IPythonKernel object at 0x00000229F684E050>>

```

```

outputs= model(**inputs)
outputs
✓ 0.3s Python
TFSequenceClassifierOutput(loss=None, logits=<tf.Tensor: shape=(1, 3), dtype=float32, numpy=array([[ -0.9634368, -2.6876621,  3.7950063]], dt

```

```
tf.nn.softmax(outputs.logits)
✓ 0.0s
<tf.Tensor: shape=(1, 3), dtype=float32, numpy=array([[0.0084931 , 0.00151441, 0.98999244]], dtype=float32)>
```

```
tf.nn.softmax(outputs.logits).numpy()
✓ 0.0s
array([[0.0084931 , 0.00151441, 0.98999244]], dtype=float32)
```

```
tf.nn.softmax(outputs.logits).numpy()[0]
✓ 0.0s
array([0.0084931 , 0.00151441, 0.98999244], dtype=float32)
```

predict_seriousness("patient died")

```
('DLT', 'Confidence: ', 0.87825364)
```

predict_seriousness("Two weeks after starting a new heart medication, I experienced severe fatigue and fainting spells. I was admitted for suspected arrhythmia. The drug was stopped, and I received a pacemaker. I stabilized after five days.")

```
('Serious', 'Confidence: ', 0.98999244)
```

Test on unseen data

```

676,Serious,"After taking allopurinol for a week, I developed fever, rash, and abnormal liver tests. Drug-induced hypersensitivity syndrome was suspected and treated in hospital."
677,D/LT,"One week after a COVID-19 vaccine, I developed Guillain-Barré syndrome with ascending weakness. I was hospitalized and treated with IVIG."
678,Non-serious,"After applying an over-the-counter eye drop, I felt brief stinging that disappeared within a minute."
679,Serious,"While on methotrexate, I developed mouth sores and low platelet count. I required hospitalization and drug discontinuation."
680,D/LT,"Post-administration of a radiotherapy sensitizer, I collapsed with hypotension and bradycardia. I was resuscitated and remained in the ICU for several days."
681,Non-serious,"After switching to a new toothpaste, I had mild gum sensitivity that resolved after a week of use."
682,Serious,"While on combination ART for HIV, I developed jaundice and elevated liver enzymes. Suspected hepatotoxicity led to drug changes and inpatient care."
683,D/LT,"After receiving an intrathecal chemotherapy, I developed paralysis and was diagnosed with spinal cord toxicity. Permanent disability remains."

```

```
import pandas as pd
```

```
# Load your test CSV (same format as training data)
```

```
test_df = pd.read_csv("test.csv") # Columns: ID, Seriousness, Narrative
```

```
test_df
```

	ID	Seriousness	Narrative
0	651	Non-serious	After applying a topical gel for muscle sorene...
1	652	Serious	Two weeks after starting lithium for mood stab...
2	653	D/LT	Shortly after receiving a chemotherapy cycle, ...
3	654	Non-serious	I took a multivitamin with iron and noticed a ...
4	655	Serious	A week after starting carbamazepine, I develop...
5	656	D/LT	Following a biologic infusion for Crohn's dise...
6	657	Non-serious	I started an anti-acne cream and experienced s...
7	658	Serious	Three days after starting linezolid for a skin...

◆ Load Model

```
import tensorflow as tf
```

```
from transformers import TFBertForSequenceClassification, BertTokenizerFast
```

```
model = TFBertForSequenceClassification.from_pretrained("./hf_models/biobert_seriousness_model")
```

```
tokenizer = BertTokenizer.from_pretrained("./hf_models/biobert_seriousness_model")
encodings = tokenizer(texts, truncation=True, padding=True, max_length=256, return_tensors="tf")
```

◆ 3. **Load and prepare test data**

```
texts = test_df["Narrative"].tolist()
labels = test_df["Seriousness"].tolist()
```

◆ 4. **Preprocess test data**

```
label_encoder = LabelEncoder()
labels_encoded = label_encoder.fit_transform(labels)
```

```
test_dataset = tf.data.Dataset.from_tensor_slices((dict(encodings), labels_encoded)).batch(16)
```

◆ 5. **Evaluate using model.evaluate()**

```
# Compile the model for evaluation
from tensorflow.keras.losses import SparseCategoricalCrossentropy
from tensorflow.keras.optimizers import Adam

model.compile(
    optimizer=Adam(learning_rate=5e-5),
```

```

    loss=SparseCategoricalCrossentropy(from_logits=True),
    metrics=["accuracy"]
)

# Evaluate
loss, accuracy = model.evaluate(test_dataset)
print(f"Test Loss: {loss:.4f}, Test Accuracy: {accuracy:.4f}")

```

```

4/4 [=====] - 7s 371ms/step - loss: 0.3199 - accuracy: 0.9200
Test Loss: 0.3199, Test Accuracy: 0.9200

```

confusion_matrix

```

from sklearn.metrics import classification_report, confusion_matrix

# Get predictions
y_pred_logits = model.predict(test_dataset).logits
y_pred = tf.argmax(y_pred_logits, axis=1).numpy()

# Decode labels (optional)
y_true = labels_encoded

print("Classification Report:\n", classification_report(y_true, y_pred))

```

```

4/4 [=====] - 5s 379ms/step
Classification Report:

```

	precision	recall	f1-score	support
0	0.83	0.94	0.88	16
1	1.00	1.00	1.00	17
2	0.93	0.82	0.88	17
accuracy			0.92	50
macro avg	0.92	0.92	0.92	50
weighted avg	0.92	0.92	0.92	50

✅ per-class recall:

Class 0 ("DLT") → Recall = 0.94

Class 1 ("Non-serious") → Recall = 1.00

Class 2 ("Serious") → Recall = 0.82

this means:

- You are **very good at catching DLTs** ✅
- You're **missing 18% of serious cases** ❌ ← needs attention